

```
In [1]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-pyth
# For example, here's several helpful packages to load
import seaborn as sns
import plotly.express as px
import matplotlib.pyplot as plt
import plotly.graph_objects as go
from tqdm import tqdm_notebook
import plotly.figure_factory as ff

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

import warnings
warnings.filterwarnings('ignore')

plt.style.use('fivethirtyeight')
%matplotlib inline
# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all file

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets preserv
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside o

In [2]: data=pd.read_csv('../input/water-potability/water_potability.csv')
data.head()
```

Out[2]:

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbon	Trihalomethanes
0	NaN	204.890455	20791.318981	7.300212	368.516441	564.308654	10.379783	86.990970
1	3.716080	129.422921	18630.057858	6.635246	NaN	592.885359	15.180013	56.329076
2	8.099124	224.236259	19909.541732	9.275884	NaN	418.606213	16.868637	66.420093
3	8.316766	214.373394	22018.417441	8.059332	356.886136	363.266516	18.436524	100.341674
4	9.092223	181.101509	17978.986339	6.546600	310.135738	398.410813	11.558279	31.997993

EDA

- ph-> pH of water
- Hardness-> Capacity of water to precipitate soap in mg/L
- Solids-> Total dissolved solids in ppm
- Chloramines-> Amount of Chloramines in ppm
- Sulfate-> Amount of Sulfates dissolved in mg/L
- Conductivity-> Electrical conductivity of water in µS/cm
- Organic_carbon-> Amount of organic carbon in ppm
- Trihalomethanes-> Amount of Trihalomethanes in µg/L
- Turbidity-> Measure of light emitting property of water in NTU (Nephelometric Turbidity Units)
- Potability-> Indicates if water is safe for human consumption

In [3]:

```
data.describe()
```

Out[3]:

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbon	Trihalomethanes
count	2785.000000	3276.000000	3276.000000	3276.000000	2495.000000	3276.000000	3276.000000	3276.000000
mean	7.080795	196.369496	22014.092526	7.122277	333.775777	426.205111	14.284970	56.329076
std	1.594320	32.879761	8768.570828	1.583085	41.416840	80.824064	3.308162	16.868637
min	0.000000	47.432000	320.942611	0.352000	129.000000	181.483754	2.200000	31.997993
25%	6.093092	176.850538	15666.690297	6.127421	307.699498	365.734414	12.065801	56.329076
50%	7.036752	196.967627	20927.833607	7.130299	333.073546	421.884968	14.218338	56.329076
75%	8.062066	216.667456	27332.762127	8.114887	359.950170	481.792304	16.557652	66.420093
max	14.000000	323.124000	61227.196008	13.127000	481.030642	753.342620	28.300000	100.341674

In [4]:

```
data.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3276 entries, 0 to 3275
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ph                    2785 non-null   float64
1   Hardness              3276 non-null   float64
2   Solids                3276 non-null   float64
3   Chloramines           3276 non-null   float64
4   Sulfate               2495 non-null   float64
5   Conductivity          3276 non-null   float64
6   Organic_carbon        3276 non-null   float64
7   Trihalomethanes       3114 non-null   float64
8   Turbidity             3276 non-null   float64
9   Potability            3276 non-null   int64
dtypes: float64(9), int64(1)
memory usage: 256.1 KB
```

```
In [5]: print('There are {} data points and {} features in the data'.format(data.shape[0],data.s
There are 3276 data points and 10 features in the data
```

Null Values

```
In [6]: sns.heatmap(data.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

```
Out[6]: <AxesSubplot:>
```



```
In [7]: for i in data.columns:
        if data[i].isnull().sum()>0:
            print("There are {} null values in {} column".format(data[i].isnull().sum(),i))
```

```
There are 491 null values in ph column
There are 781 null values in Sulfate column
There are 162 null values in Trihalomethanes column
```

Handling Null Values

PH

```
In [8]: data['ph'].describe()
```

```
Out[8]: count      2785.000000  
mean         7.080795  
std          1.594320  
min           0.000000  
25%          6.093092  
50%          7.036752  
75%          8.062066  
max          14.000000  
Name: ph, dtype: float64
```


Filling the missing values by mean

```
In [9]: data['ph_mean']=data['ph'].fillna(data['ph'].mean())
```

```
In [10]: data['ph_mean'].isnull().sum()
```

```
Out[10]: 0
```

```
In [11]: fig = plt.figure()  
ax = fig.add_subplot(111)  
data['ph'].plot(kind='kde', ax=ax)  
data.ph_mean.plot(kind='kde', ax=ax, color='red')  
lines, labels = ax.get_legend_handles_labels()  
ax.legend(lines, labels, loc='best')  
plt.show()
```



The distribution is not uniform

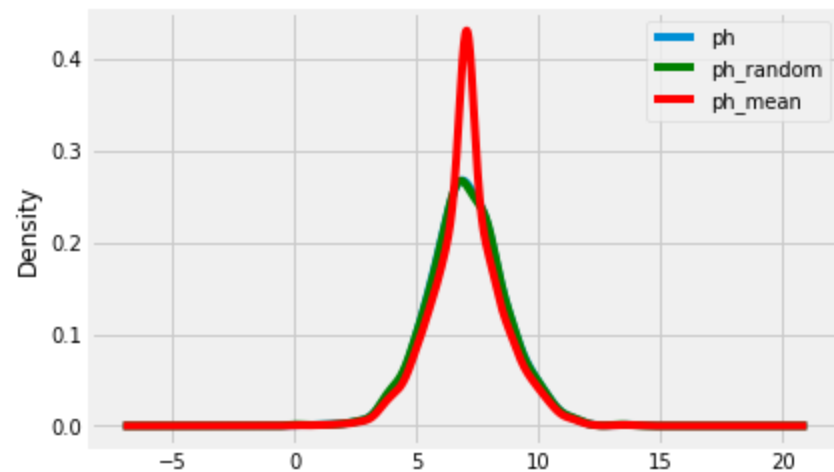
Filling the data with random values

```
In [12]: def impute_nan(df,variable):  
df[variable+"_random"]=df[variable]  
##It will have the random sample to fill the na  
random_sample=df[variable].dropna().sample(df[variable].isnull().sum(),random_state=  
##pandas need to have same index in order to merge the dataset  
random_sample.index=df[df[variable].isnull()].index  
df.loc[df[variable].isnull(),variable+'_random']=random_sample
```

```
In [13]: impute_nan(data,"ph")
```

```
In [14]: fig = plt.figure()  
ax =fig.add_subplot(111)  
data['ph'].plot(kind='kde', ax=ax)
```


```
data.ph_random.plot(kind='kde', ax=ax, color='green')
data.ph_mean.plot(kind='kde', ax=ax, color='red')
lines, labels = ax.get_legend_handles_labels()
ax.legend(lines, labels, loc='best')
plt.show()
```



Uniform distribution with random initialization

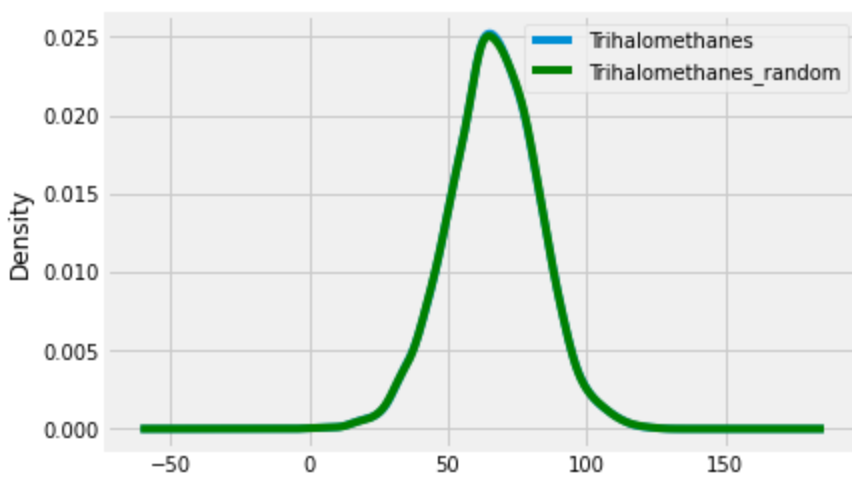
```
In [15]: impute_nan(data, "Sulfate")
```

```
In [16]: fig = plt.figure()
ax = fig.add_subplot(111)
data['Sulfate'].plot(kind='kde', ax=ax)
data["Sulfate_random"].plot(kind='kde', ax=ax, color='green')
#data.ph_mean.plot(kind='kde', ax=ax, color='red')
lines, labels = ax.get_legend_handles_labels()
ax.legend(lines, labels, loc='best')
plt.show()
```



```
In [17]: impute_nan(data, "Trihalomethanes")
```

```
In [18]: fig = plt.figure()
ax = fig.add_subplot(111)
data['Trihalomethanes'].plot(kind='kde', ax=ax)
data.Trihalomethanes_random.plot(kind='kde', ax=ax, color='green')
lines, labels = ax.get_legend_handles_labels()
ax.legend(lines, labels, loc='best')
plt.show()
```




```
In [19]: data=data.drop(['ph','Sulfate','Trihalomethanes','ph_mean'],axis=1)
```

```
In [20]: data.isnull().sum()
```

```
Out[20]: Hardness          0
         Solids           0
         Chloramines       0
         Conductivity      0
         Organic_carbon    0
         Turbidity         0
         Potability        0
         ph_random         0
         Sulfate_random    0
         Trihalomethanes_random  0
         dtype: int64
```

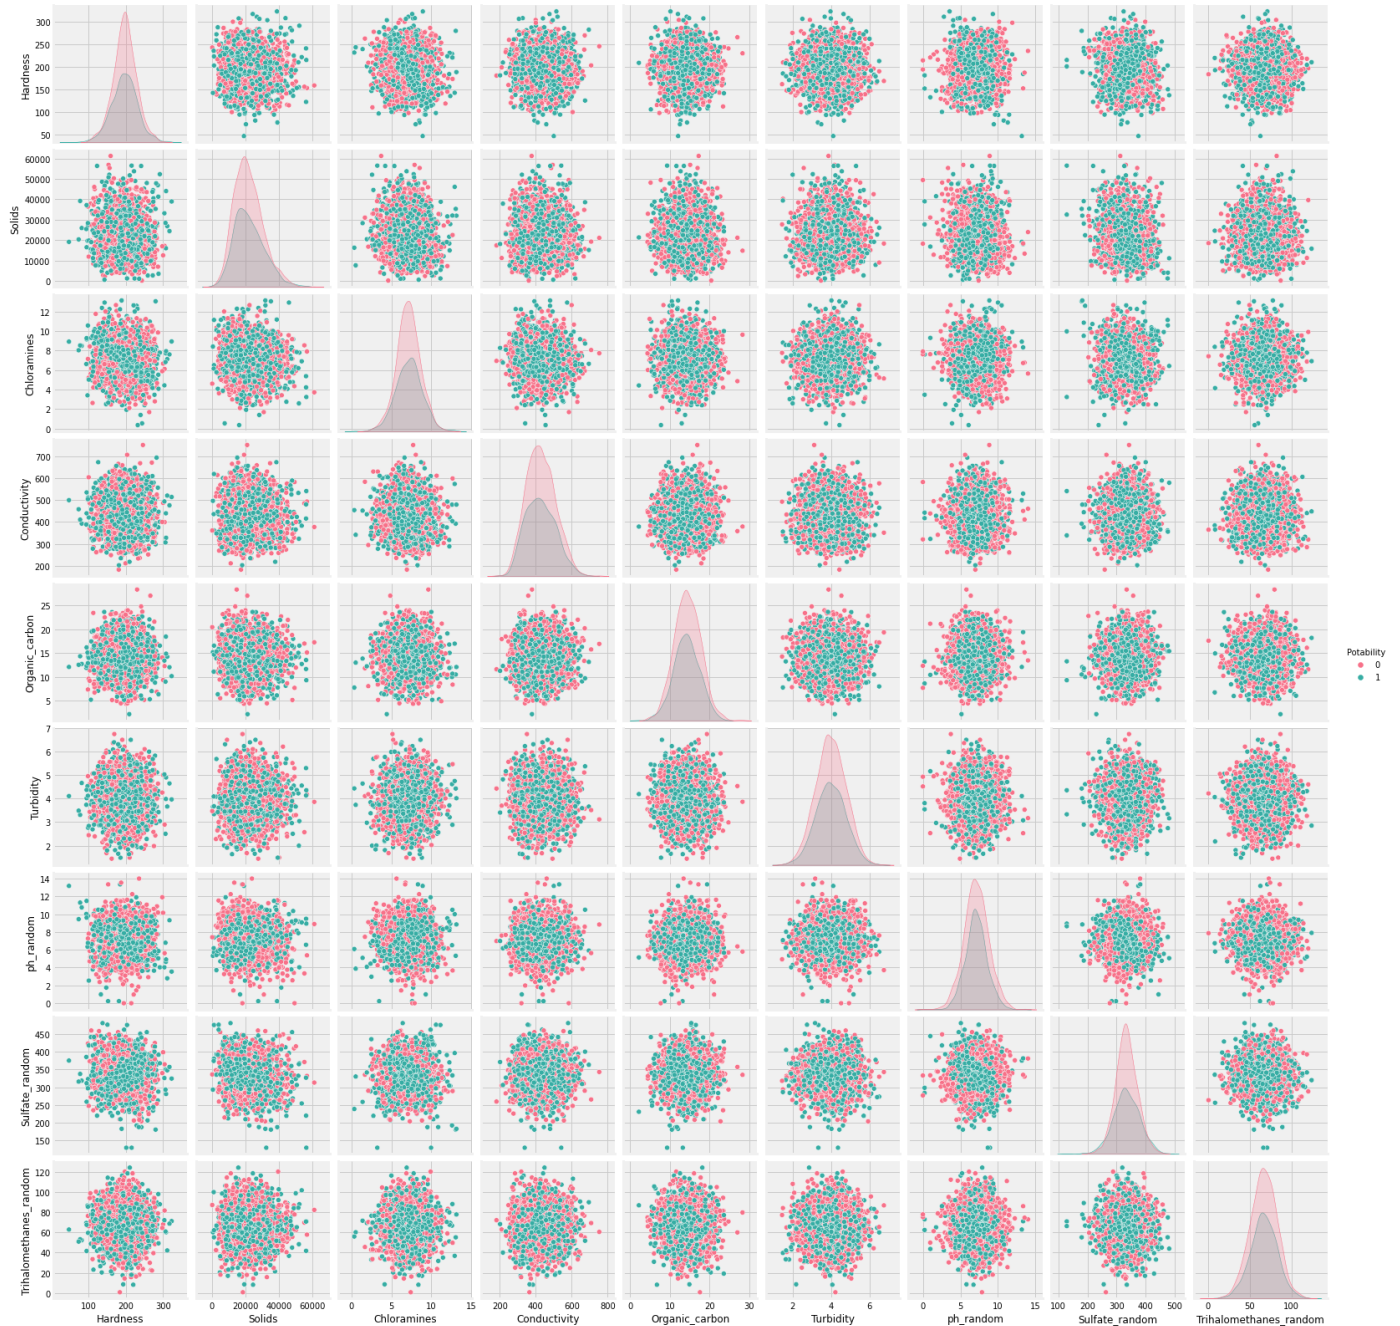
Check for Correlation

```
In [21]: plt.figure(figsize=(20, 17))
         matrix = np.triu(data.corr())
         sns.heatmap(data.corr(), annot=True,linewidth=.8, mask=matrix, cmap="rocket",cbar=False)
```




There are no correlated columns presebt in the data

```
In [22]: sns.pairplot(data, hue="Potability", palette="husl");
```



```
In [23]: def distributionPlot(data):
    """
    Creates distribution plot.
    """
    fig = plt.figure(figsize=(20, 20))
    for i in tqdm_notebook(range(0, len(data.columns))):
        fig.add_subplot(np.ceil(len(data.columns)/3), 3, i+1)
        sns.distplot(
            data.iloc[:, i], color="lightcoral", rug=True)
        fig.tight_layout(pad=3.0)
    plot_data = data.drop(['Potability'], axis =1)
    distributionPlot(plot_data)
```

0%| | 0/9 [00:00<?, ?it/s]



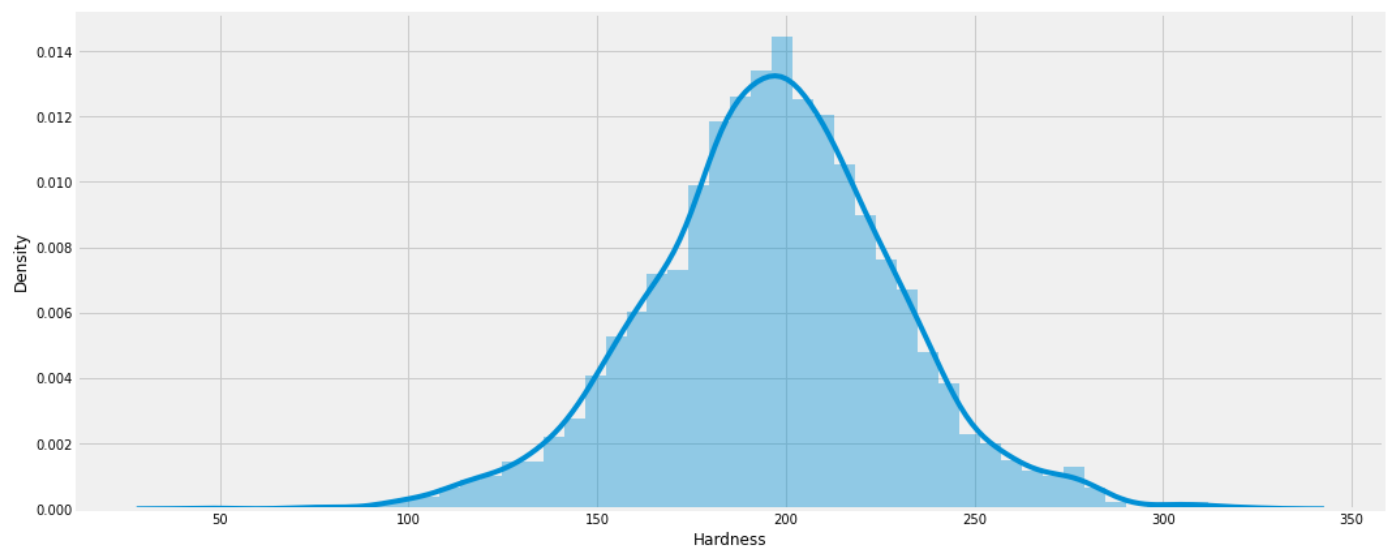
Hardness

```
In [24]: data['Hardness'].describe()
```


```
Out[24]: count    3276.000000
mean      196.369496
std       32.879761
min       47.432000
25%      176.850538
50%      196.967627
75%      216.667456
max       323.124000
Name: Hardness, dtype: float64
```

```
In [25]: plt.figure(figsize = (16, 7))
sns.distplot(data['Hardness'])
plt.title('Distribution Plot of Hardness\n', fontsize = 20)
plt.show()
```


Distribution Plot of Hardness



```
In [26]: # basic scatter plot
fig = px.scatter(data, range(data['Hardness'].count()), sorted(data['Hardness']),
                 color=data['Potability'],
                 labels={
                     'x': "Count",
                     'y': "Hardness",
                     'color': 'Potability'
                 }, template = 'plotly_dark')
fig.update_layout(title='Hardness wrt Potability')
fig.show()
```



```
In [27]: px.histogram(data_frame = data, x = 'Hardness', nbins = 10, color = 'Potability', margin
              template = 'plotly_dark')
```




Solids

```
In [28]: data['Solids'].describe()
```


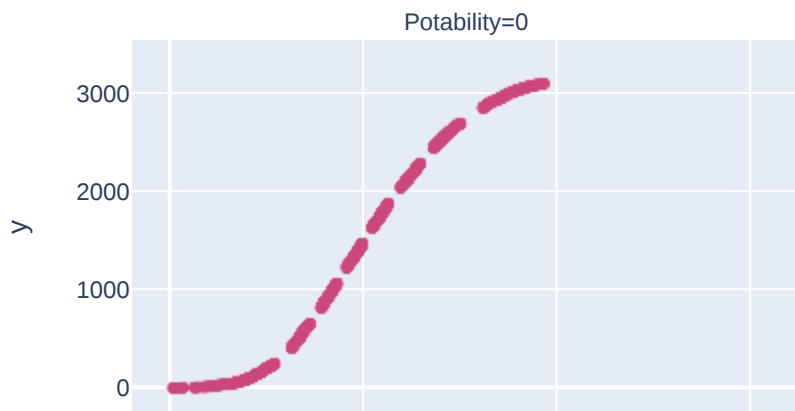
```
Out[28]: count      3276.000000  
mean      22014.092526  
std       8768.570828  
min       320.942611  
25%      15666.690297  
50%      20927.833607  
75%      27332.762127  
max       61227.196008  
Name: Solids, dtype: float64
```

```
In [29]: plt.figure(figsize = (16, 7))  
sns.distplot(data['Solids'])  
plt.title('Distribution Plot of Solids\n', fontsize = 20)  
plt.show()
```

Distribution Plot of Solids



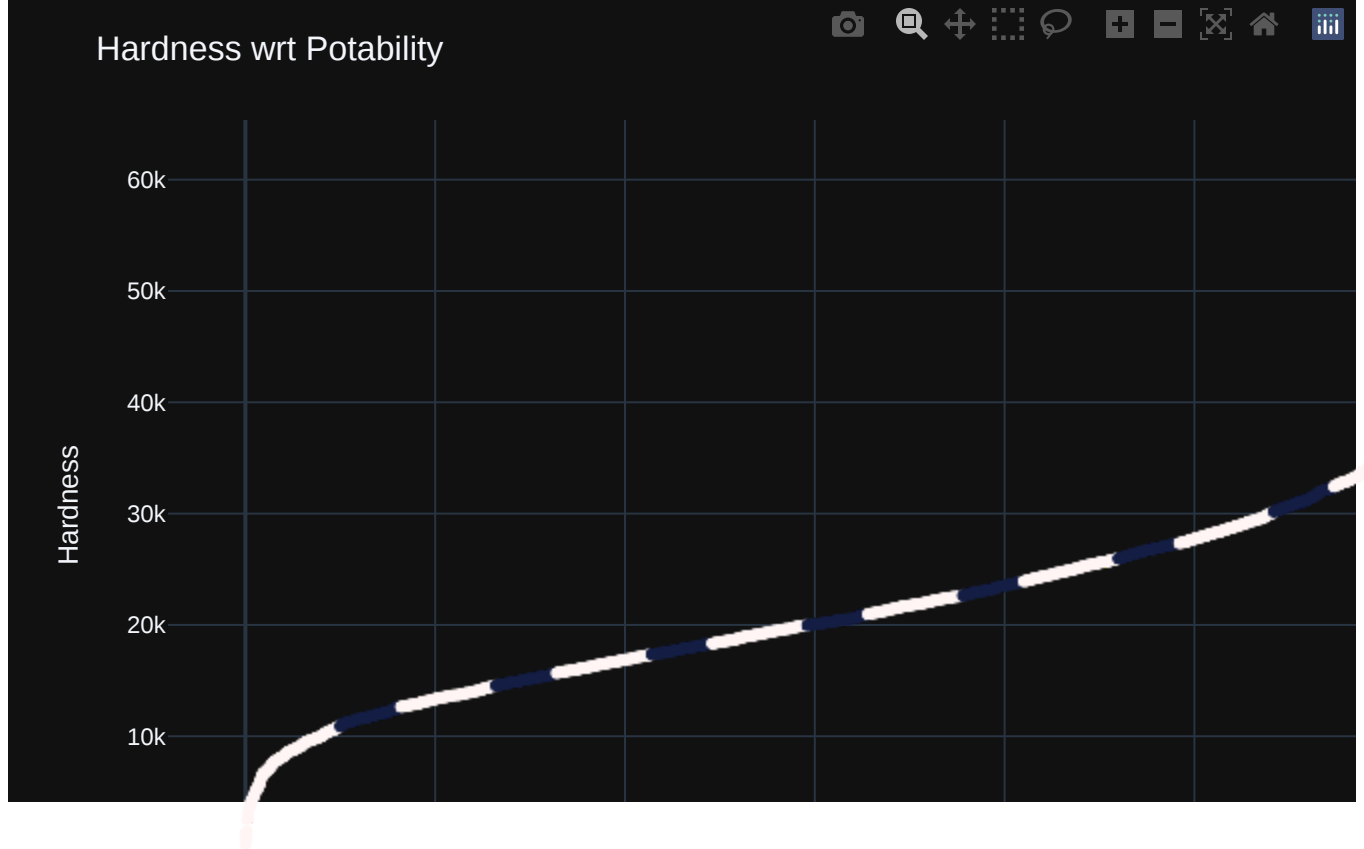
```
In [30]: fig = px.scatter(data, sorted(data["Solids"]), range(data["Solids"].count()), color="Potability",
                        facet_row="Potability")
fig.show()
```



```
In [31]: px.histogram(data_frame = data, x = 'Solids', nbins = 10, color = 'Potability', marginal
                    template = 'plotly_dark')
```



```
In [32]: # basic scatter plot
fig = px.scatter(data, range(data['Solids'].count()), sorted(data['Solids']),
                 color=data['Potability'],
                 labels={
                     'x': "Count",
                     'y': "Hardness",
                     'color': 'Potability'
                 },
                 color_continuous_scale=px.colors.sequential.tempo,
                 template = 'plotly_dark')
fig.update_layout(title='Hardness wrt Potability')
fig.show()
```




Chloramines

```
In [33]: data['Chloramines'].describe()
```

```
Out[33]: count    3276.000000
mean         7.122277
std          1.583085
min           0.352000
25%          6.127421
50%          7.130299
75%          8.114887
max          13.127000
Name: Chloramines, dtype: float64
```


```
In [34]: plt.figure(figsize = (16, 7))
sns.distplot(data['Chloramines'])
plt.title('Distribution Plot of Chloramines\n', fontsize = 20)
plt.show()
```

Distribution Plot of Chloramines



```
In [35]: fig = px.line(x=range(data['Chloramines'].count()), y=sorted(data['Chloramines']), color=
            'x': "Count",
            'y': "Chloramines",
            'color': 'Potability'


            }, template = 'plotly_dark')
fig.update_layout(title='Chloramines wrt Potability')
fig.show()
```



```
In [36]: fig = px.box(x = 'Chloramines', data_frame = data, template = 'plotly_dark')
fig.update_layout(title='Chloramines')
```



```
fig.show()
```



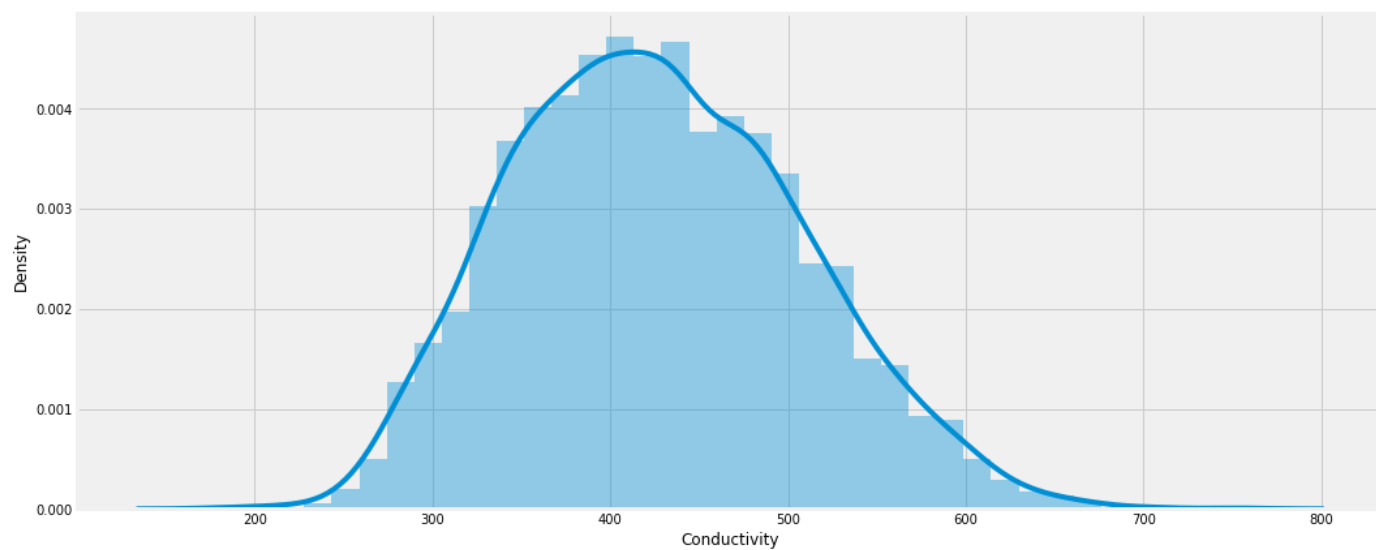
Conductivity

```
In [37]: data["Conductivity"].describe()
```

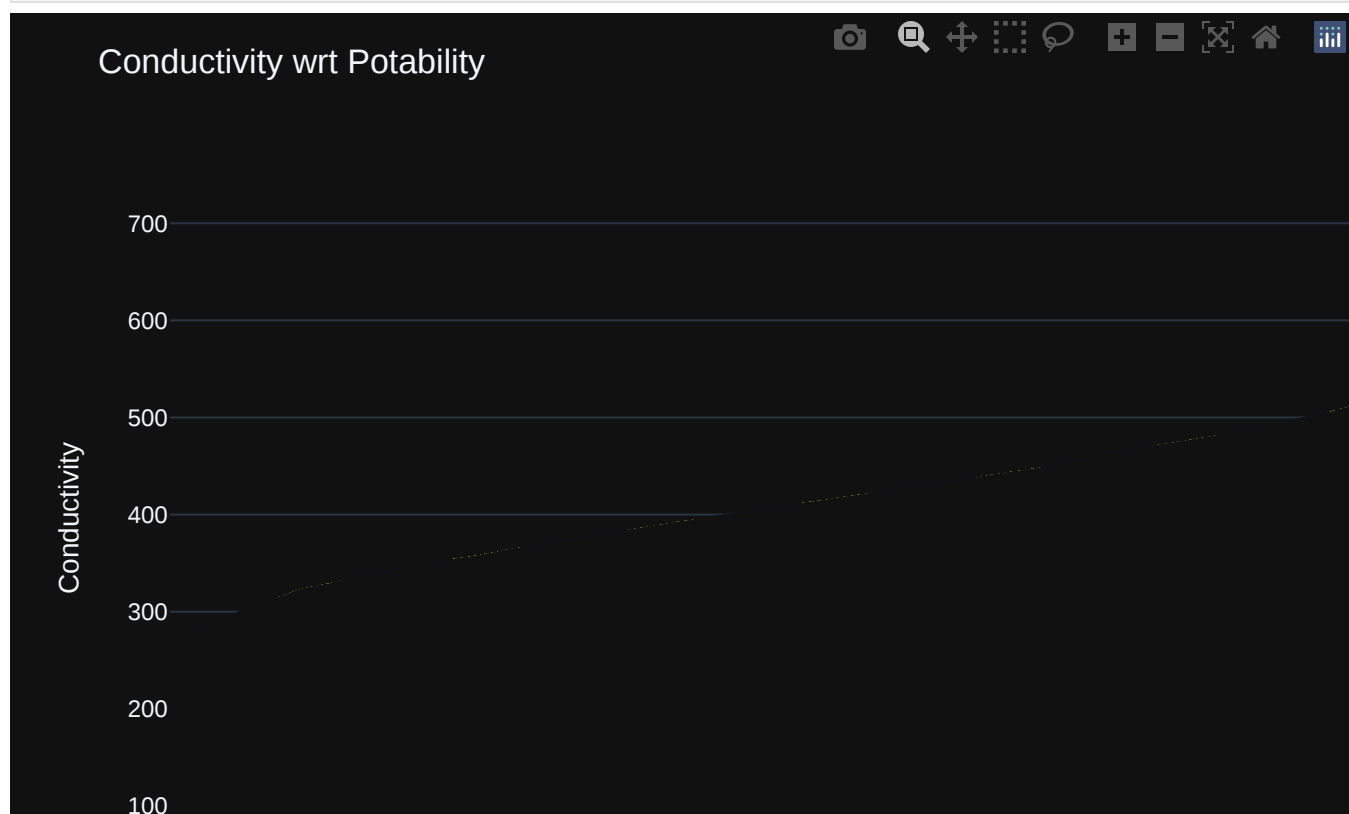
```
Out[37]: count    3276.000000  
mean      426.205111  
std       80.824064  
min       181.483754  
25%       365.734414  
50%       421.884968  
75%       481.792304  
max       753.342620  
Name: Conductivity, dtype: float64
```

```
In [38]: plt.figure(figsize = (16, 7))  
sns.distplot(data['Conductivity'])  
plt.title('Distribution Plot of Conductivity\n', fontsize = 20)  
plt.show()
```

Distribution Plot of Conductivity




```
In [39]: fig = px.bar(data, x=range(data['Conductivity'].count()),
                    y=sorted(data['Conductivity']), labels={
                        'x': "Count",
                        'y': "Conductivity",
                        'color': 'Potability'
                    },
                    color=data['Potability'],
                    template = 'plotly_dark')
fig.update_layout(title='Conductivity wrt Potability')
fig.show()
```



```
In [40]: group_labels = ['distplot'] # name of the dataset

fig = ff.create_distplot([data['Conductivity']], group_labels)
fig.show()
```




Organic_carbon

```
In [41]: data['Organic_carbon'].describe()
```

```
Out[41]: count      3276.000000
mean        14.284970
std         3.308162
min         2.200000
25%        12.065801
50%        14.218338
75%        16.557652
max         28.300000
Name: Organic_carbon, dtype: float64
```

```
In [42]: group_labels = ['Organic_carbon'] # name of the dataset

fig = ff.create_distplot([data['Organic_carbon']], group_labels)
fig.show()
```



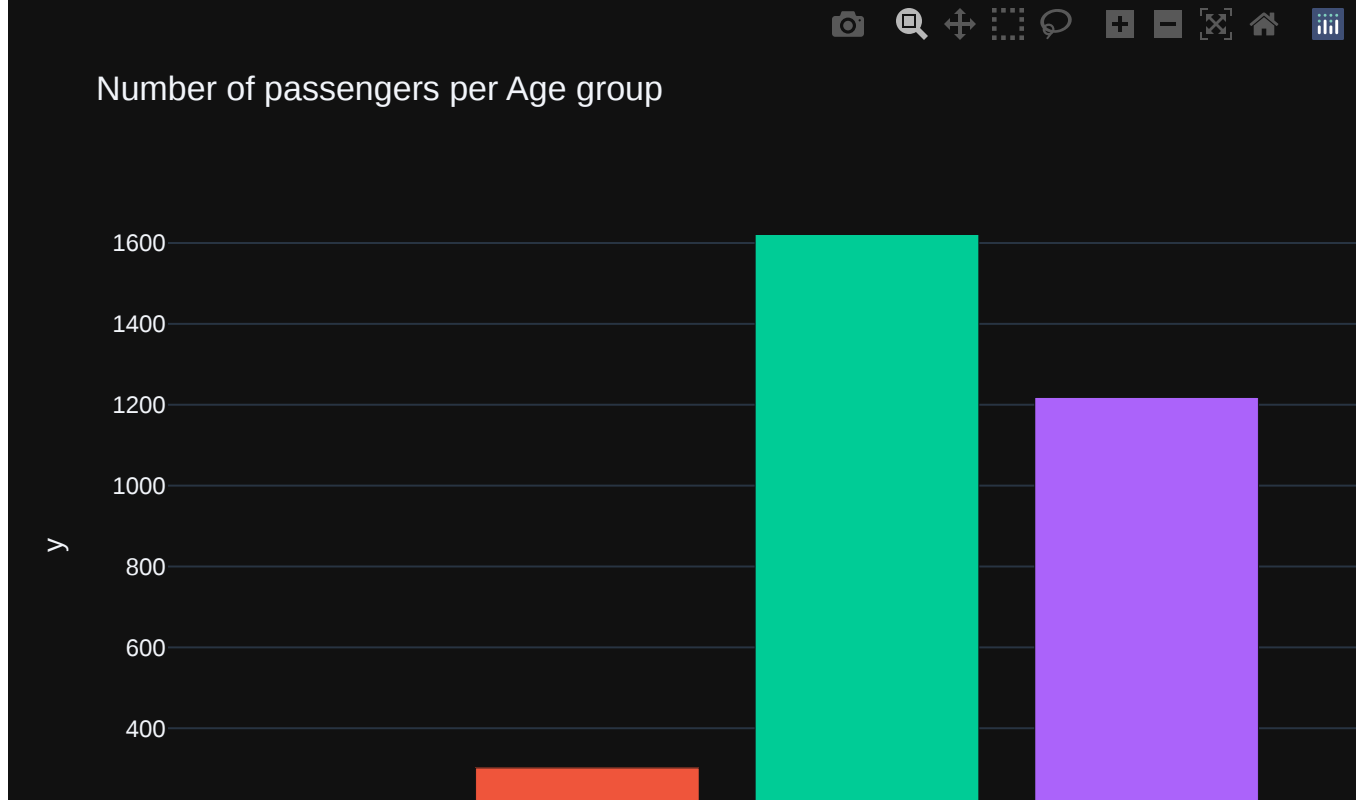
```

In [43]: dt_5=data[data['Organic_carbon']<5]
dt_5_10=data[(data['Organic_carbon']>5)&(data['Organic_carbon']<10)]
dt_10_15=data[(data['Organic_carbon']>10)&(data['Organic_carbon']<15)]
dt_15_20=data[(data['Organic_carbon']>15)&(data['Organic_carbon']<20)]
dt_20_25=data[(data['Organic_carbon']>20)&(data['Organic_carbon']<25)]
dt_25=data[(data['Organic_carbon']>25)]

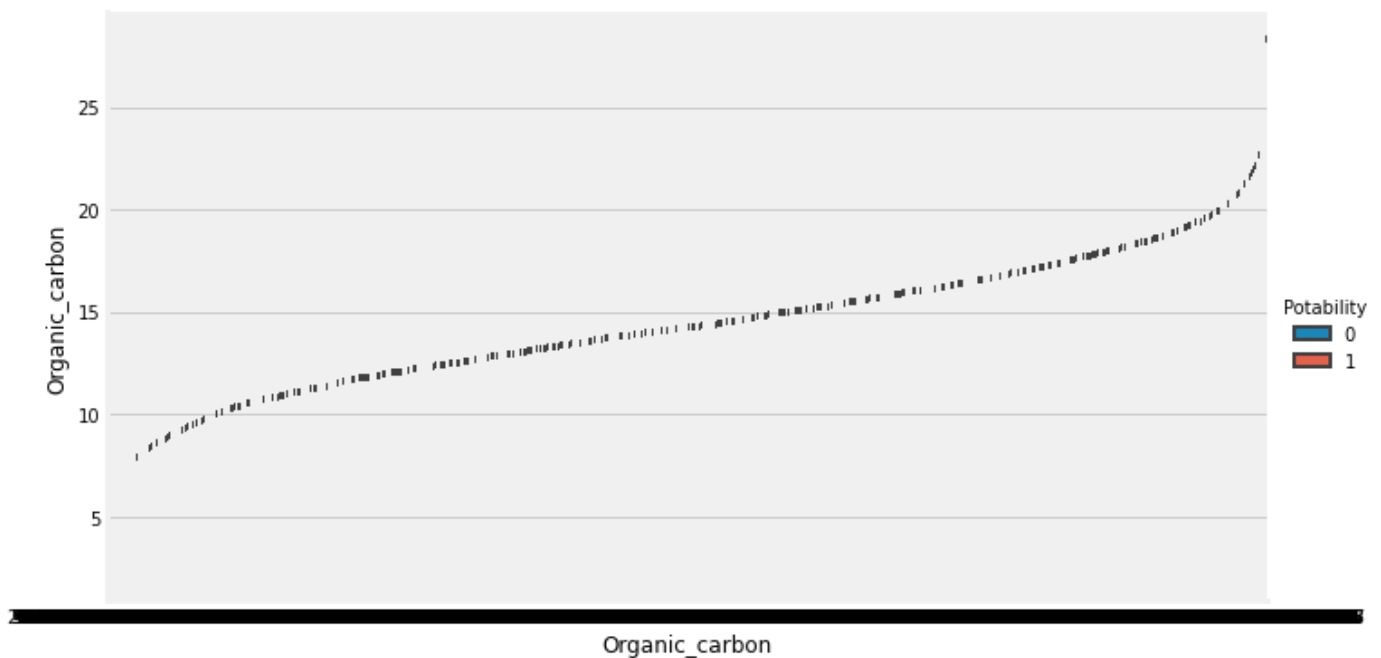
x_Age = ['5', '5-10', '10-15', '15-20', '25+']
y_Age = [len(dt_5.values), len(dt_5_10.values), len(dt_10_15.values), len(dt_15_20.values),
len(dt_25.values)]

px.bar(data_frame = data, x = x_Age, y = y_Age, color = x_Age, template = 'plotly_dark',
title = 'Number of passengers per Age group')

```



```
In [44]: sns.catplot(x = 'Organic_carbon', y = 'Organic_carbon', hue = 'Potability', data = data,
                    height = 5, aspect = 2)
plt.show()
```




Turbidity

```
In [45]: data['Turbidity'].describe()
```

```
Out[45]: count      3276.000000  
         mean        3.966786  
         std         0.780382  
         min         1.450000  
         25%         3.439711  
         50%         3.955028  
         75%         4.500320  
         max         6.739000  
         Name: Turbidity, dtype: float64
```

```
In [46]: group_labels = ['Turbidity'] # name of the dataset  
  
fig = ff.create_distplot([data['Turbidity']], group_labels)  
fig.show()
```

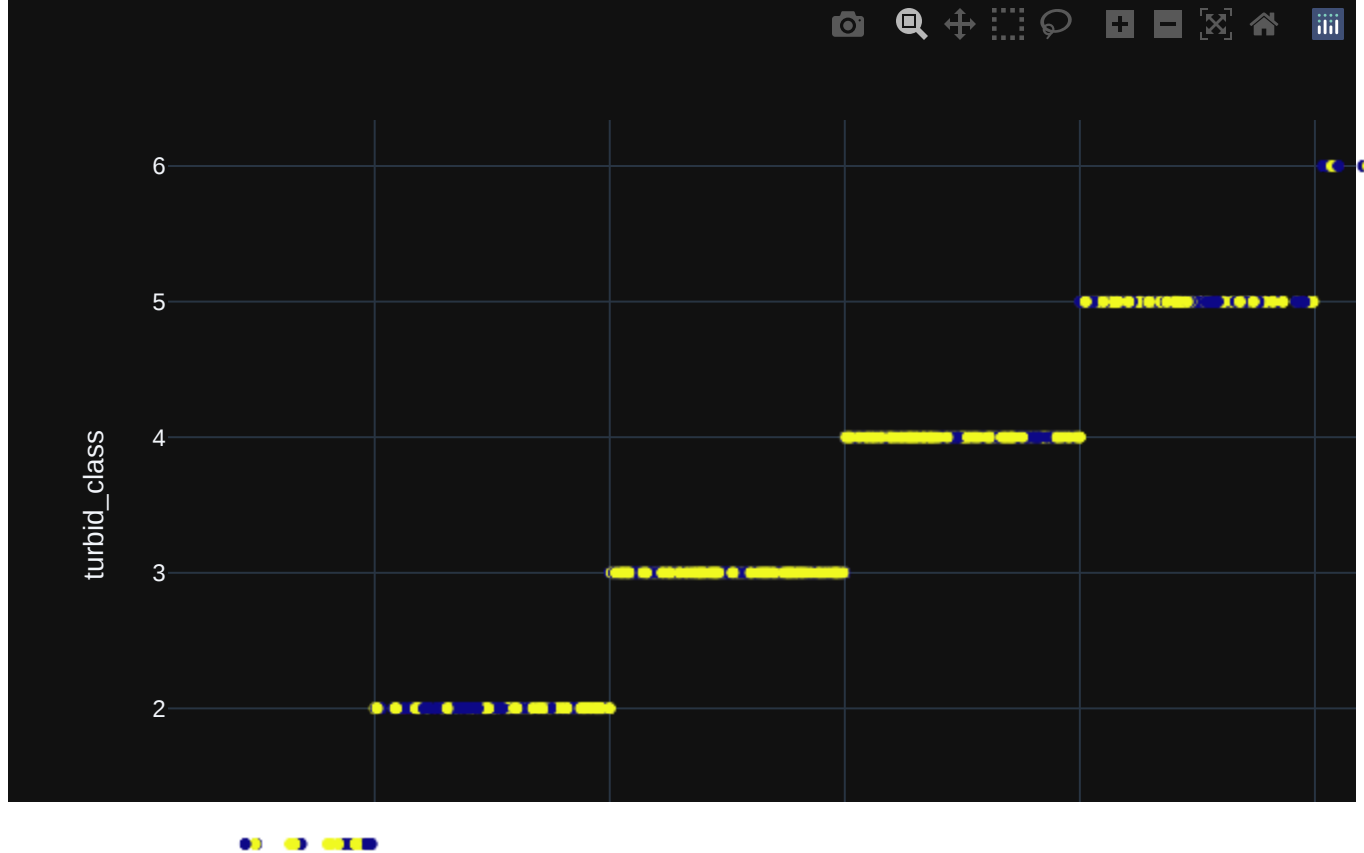


```
In [47]: data['turbid_class']=data['Turbidity'].astype(int)
```

```
In [48]: data['turbid_class'].unique()
```


```
Out[48]: array([2, 4, 3, 5, 6, 1])
```

```
In [49]: px.scatter(data_frame = data, x = 'Turbidity', y = 'turbid_class', color = 'Potability',
```



```
In [50]: fig = px.pie(data,
                    values=data['turbid_class'].value_counts(),
                    names=data['turbid_class'].value_counts().keys(),
                    )
fig.update_layout(
    title='turbid_class',
    template = 'plotly_dark'
)
fig.show()
```

turbid_class



```
In [51]: data=data.drop(['turbid_class'],axis=1)
```

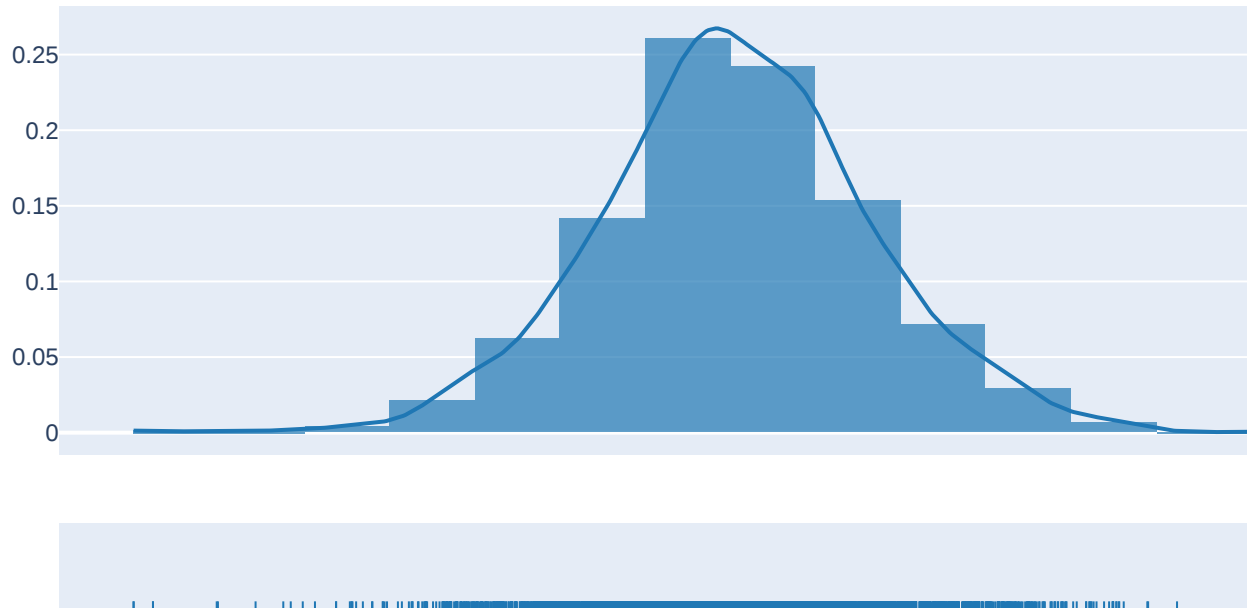
ph_random

```
In [52]: data['ph_random'].describe()
```


```
Out[52]: count    3276.000000
mean        7.071639
std         1.607991
min         0.000000
25%         6.081460
50%         7.029490
75%         8.063147
max         14.000000
Name: ph_random, dtype: float64
```

```
In [53]: group_labels = ['ph_random'] # name of the dataset

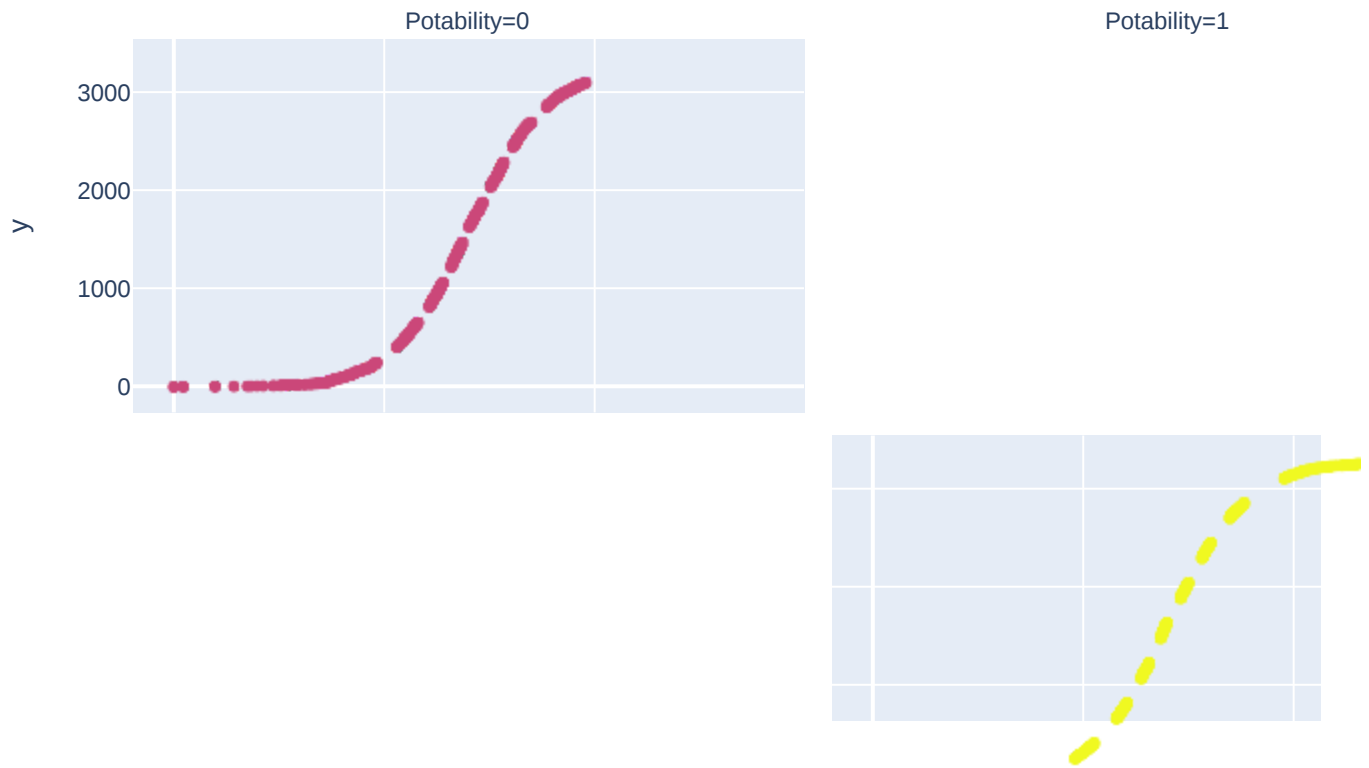
fig = ff.create_distplot([data['ph_random']], group_labels)
fig.show()
```

```
In [54]: px.histogram(data_frame = data, x = 'ph_random', nbins = 10, color = 'Potability', margin=10,
                    template = 'plotly_dark')
```



```
In [55]: fig = px.scatter(data, sorted(data["ph_random"]), range(data["ph_random"].count()), color=
          facet_row="Potability")
fig.show()
```




Sulfate_random

```
In [56]: data['Sulfate_random'].describe()
```

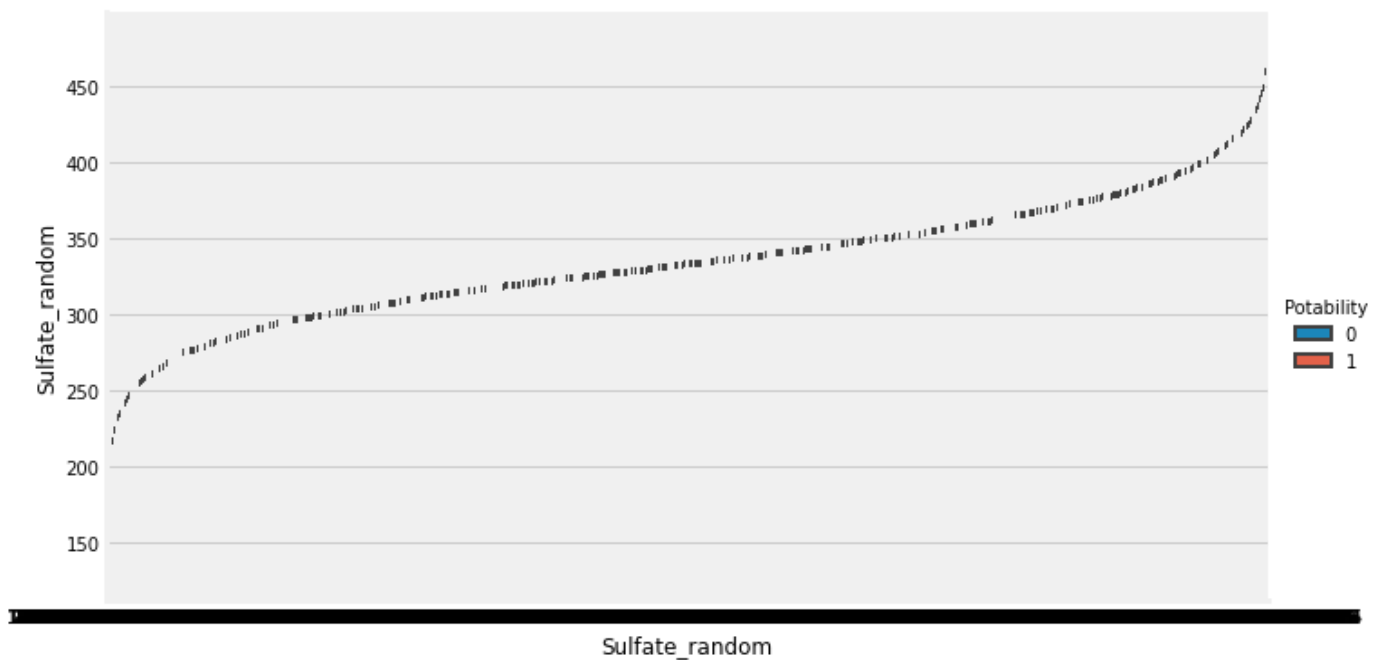
```
Out[56]: count    3276.000000
mean      333.430954
std       41.026947
min       129.000000
25%       307.523159
50%       332.879578
75%       359.710517
max       481.030642
Name: Sulfate_random, dtype: float64
```

```
In [57]: group_labels = ['distplot'] # name of the dataset

fig = ff.create_distplot([data['Sulfate_random']], group_labels)
fig.show()
```



```
In [58]: sns.catplot(x = 'Sulfate_random', y = 'Sulfate_random', hue = 'Potability', data = data,
                    height = 5, aspect = 2)
plt.show()
```




Trihalomethanes_random

```
In [59]: data['Trihalomethanes_random'].describe()
```


```
Out[59]: count    3276.000000
mean       66.419200
std        16.184832
min         0.738000
25%        55.861675
50%        66.639068
75%        77.384166
max       124.000000
Name: Trihalomethanes_random, dtype: float64
```

```
In [60]: group_labels = ['Trihalomethanes_random'] # name of the dataset

fig = ff.create_distplot([data['Trihalomethanes_random']], group_labels)
fig.show()
```




```
In [61]: fig = px.box(x = 'Trihalomethanes_random', data_frame = data, template = 'plotly_dark')
fig.update_layout(title='Trihalomethanes_random')
fig.show()
```



```
In [62]: fig = px.line(x=range(data['Trihalomethanes_random'].count()), y=sorted(data['Trihalomet
        'x': "Count",
        'y': "Trihalomethanes",
        'color': 'Potability'

        }, template = 'plotly_dark')
fig.update_layout(title='Trihalomethane wrt Potability')
fig.show()
```




Potability


```
In [63]: data['Potability'].describe()
```

```
Out[63]: count    3276.000000
mean         0.390110
std          0.487849
min          0.000000
25%          0.000000
50%          0.000000
75%          1.000000
max          1.000000
Name: Potability, dtype: float64
```

```
In [64]: px.histogram(data_frame = data, x = 'Potability', color = 'Potability', marginal = 'box'
                    template = 'plotly_dark')
```



```
In [65]: fig = px.pie(data,
                    values=data['Potability'].value_counts(),
                    names=data['Potability'].value_counts().keys(),
                    )
fig.update_layout(
    title='Potability',
    template = 'plotly_dark'
)
fig.show()
```

Data Preprocessing

```
In [66]: from sklearn.preprocessing import StandardScaler  
from sklearn.model_selection import train_test_split
```

```
In [67]: X=data.drop(['Potability'],axis=1)  
y=data['Potability']
```

Since the data is not in a uniform shape, we scale the data using standard scalar

```
In [68]: scaler = StandardScaler()  
x=scaler.fit_transform(X)
```

```
In [69]: # split the data to train and test set  
x_train,x_test,y_train,y_test = train_test_split(x,y,train_size=0.85,random_state=42)
```

```
print("training data shape:-{} labels{}".format(x_train.shape,y_train.shape))  
print("testing data shape:-{} labels{}".format(x_test.shape,y_test.shape))
```

```
training data shape:-(2784, 9) labels(2784,)  
testing data shape:-(492, 9) labels(492,)
```


Modeling

Logistic Regression

```
In [70]: from sklearn.linear_model import LogisticRegression
log = LogisticRegression(random_state=0).fit(x_train, y_train)
log.score(x_test, y_test)
```

Out[70]: 0.6219512195121951

```
In [71]: # Confusion matrix
from sklearn.metrics import confusion_matrix
# Make Predictions
pred1=log.predict(np.array(x_test))
plt.title("Confusion Matrix testing data")
sns.heatmap(confusion_matrix(y_test, pred1), annot=True, cbar=False)
plt.legend()
plt.show()
```




K Nearest Neighbours

```
In [72]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [73]: knn = KNeighborsClassifier(n_neighbors=2)
# Train the model using the training sets
knn.fit(x_train, y_train)

#Predict Output
predicted= knn.predict(x_test) # 0:Overcast, 2:Mild
```

```
In [74]: # Confusion matrix
from sklearn.metrics import confusion_matrix
# Make Predictions
pred1=knn.predict(np.array(x_test))
plt.title("Confusion Matrix testing data")
sns.heatmap(confusion_matrix(y_test, pred1), annot=True, cbar=False)
plt.legend()
plt.show()
```



SVM


```
In [75]: from sklearn import svm
from sklearn.metrics import accuracy_score
```

```
In [76]: svmc = svm.SVC()
svmc.fit(x_train, y_train)

y_pred = svmc.predict(x_test)
print(accuracy_score(y_test, y_pred))
```

0.6808943089430894

```
In [77]: # Confusion matrix
from sklearn.metrics import confusion_matrix
# Make Predictions
pred1=svmc.predict(np.array(x_test))
plt.title("Confusion Matrix testing data")
sns.heatmap(confusion_matrix(y_test, pred1), annot=True, cbar=False)
plt.legend()
plt.show()
```



Decision Tree


```
In [78]: from sklearn import tree
from sklearn.metrics import accuracy_score
```

```
In [79]: tre = tree.DecisionTreeClassifier()
tre = tre.fit(x_train, y_train)

y_pred = tre.predict(x_test)
print(accuracy_score(y_test, y_pred))
```

0.5487804878048781

```
In [80]: # Confusion matrix
from sklearn.metrics import confusion_matrix
# Make Predictions
pred1=tre.predict(np.array(x_test))
plt.title("Confusion Matrix testing data")
sns.heatmap(confusion_matrix(y_test, pred1), annot=True, cbar=False)
plt.legend()
plt.show()
```



Random Forest

```
In [81]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
```


```
In [82]: # create the model
model_rf = RandomForestClassifier(n_estimators=500, oob_score=True, random_state=100)

# fitting the model
model_rf=model_rf.fit(x_train, y_train)

y_pred = model_rf.predict(x_test)
print(accuracy_score(y_test, y_pred))
```

0.6788617886178862

```
In [83]: # Confusion matrix
from sklearn.metrics import confusion_matrix
# Make Predictions
pred1=model_rf.predict(np.array(x_test))
plt.title("Confusion Matrix testing data")
sns.heatmap(confusion_matrix(y_test, pred1), annot=True, cbar=False)
plt.legend()
plt.show()
```



XG Boost

```
In [84]: from xgboost import XGBClassifier
from sklearn.metrics import r2_score
```

```
xgb = XGBClassifier(colsample_bylevel= 0.9,
                    colsample_bytree = 0.8,
                    gamma=0.99,
                    max_depth= 5,
                    min_child_weight= 1,
                    n_estimators= 8,
                    nthread= 5,
                    random_state= 0,
                    )
xgb.fit(x_train,y_train)
```

[14:31:40] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

```
Out[84]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=0.9,
                    colsample_bynode=1, colsample_bytree=0.8, gamma=0.99, gpu_id=-1,
                    importance_type='gain', interaction_constraints='',
                    learning_rate=0.300000012, max_delta_step=0, max_depth=5,
                    min_child_weight=1, missing=nan, monotone_constraints='()',
                    n_estimators=8, n_jobs=5, nthread=5, num_parallel_tree=1,
                    random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
                    subsample=1, tree_method='exact', validate_parameters=1,
                    verbosity=None)
```


```
In [85]: print('Accuracy of XGBoost classifier on training set: {:.2f}'
            .format(xgb.score(x_train, y_train)))
print('Accuracy of XGBoost classifier on test set: {:.2f}'
      .format(xgb.score(x_test, y_test)))
```

Accuracy of XGBoost classifier on training set: 0.72
Accuracy of XGBoost classifier on test set: 0.63

```
In [86]: from sklearn.metrics import confusion_matrix

conf_matrix = confusion_matrix(y_true=y_test, y_pred=y_pred)
plt.figure(figsize = (15, 8))
sns.set(font_scale=1.4) # for label size
sns.heatmap(conf_matrix, annot=True, annot_kws={"size": 16}, cbar=False, linewidths = 1)
plt.title("Test Confusion Matrix")
plt.xlabel("Predicted class")
plt.ylabel("Actual class")
```

```
plt.savefig('conf_test.png')
plt.show()
```



SVM tuned

```
In [87]: from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
svc=SVC()
param_grid={'C':[1.2,1.5,2.2,3.5,3.2,4.1], 'kernel':['linear', 'poly', 'rbf', 'sigmoid'],
gridsearch=GridSearchCV(svc,param_grid=param_grid,n_jobs=-1,verbose=4,cv=3)
gridsearch.fit(x_train,y_train)
```

Fitting 3 folds for each of 240 candidates, totalling 720 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 17 tasks      | elapsed: 2.9s
[Parallel(n_jobs=-1)]: Done 90 tasks      | elapsed: 7.5s
[Parallel(n_jobs=-1)]: Done 213 tasks     | elapsed: 15.3s
[Parallel(n_jobs=-1)]: Done 384 tasks     | elapsed: 26.8s
[Parallel(n_jobs=-1)]: Done 605 tasks     | elapsed: 42.5s
[Parallel(n_jobs=-1)]: Done 720 out of 720 | elapsed: 51.0s finished
```

```
Out[87]: GridSearchCV(cv=3, estimator=SVC(), n_jobs=-1,
param_grid={'C': [1.2, 1.5, 2.2, 3.5, 3.2, 4.1],
'degree': [1, 2, 4, 8, 10], 'gamma': ['scale', 'auto'],
'kernel': ['linear', 'poly', 'rbf', 'sigmoid']},
verbose=4)
```

```
In [88]: y_pred=gridsearch.predict(x_test)
from sklearn.metrics import confusion_matrix

conf_matrix = confusion_matrix(y_true=y_test, y_pred=y_pred)
plt.figure(figsize = (15, 8))
sns.set(font_scale=1.4) # for label size
sns.heatmap(conf_matrix, annot=True, annot_kws={"size": 16},cbar=False, linewidths = 1)
plt.title("Test Confusion Matrix")
plt.xlabel("Predicted class")
plt.ylabel("Actual class")
plt.savefig('conf_test.png')
plt.show()
```

